

# Predicting Housing Valuations In a Volatile Economy

## 1 Description

In this project we analyse and predict housing value in a volatile market over a four years window. The dataset is from kaggle.com, including the characteristics of sold houses and the microeconomics indexes. While cleaning the data, we use ggplot to plot variables, making 24 graphs (including one interactive plot). We use Multivariate Imputation by Chained Equations (mice), for missing variables imputation. Finally we run a random search XGBoost with 1000 draws to find the best model, which outperforms simple regression by about 50 percent.

```
# == Data Visualisation and Wrangling == #  
library(tidyverse)  
library(data.table)  
library(lubridate)  
library(ggthemes)  
  
# == Imputing Missing Data == #  
library(mice)  
library(lattice)  
  
# == Interactive Time series == #  
library(dygraphs)  
library(xts)  
  
# == XGBoost == #  
library(xgboost)  
library(Metrics)
```

### 1.0.0.1 set seed

```
set.seed(1234)
```

## 2 Loading data and initial preparation

### 2.0.0.1 set seed

```
set.seed(1234)
```

```
df = read.csv("data.csv" , header= TRUE)  
macro = read.csv("macro.csv" , header= TRUE)
```

## 3 checking the data

## The data dimensions

```
dim(df)
```

```
## [1] 30471 292
```

Converting data columns to appropriate format.

```
df$timestamp <- as.Date(df$timestamp)
macro$timestamp <- as.Date(macro$timestamp)
```

We also limit the number of variables/columns as this project is a demonstration and the resources (time/computation) are limited for intended analysis.

```
df <- df %>% select(timestamp, full_sq, life_sq, floor,
                  max_floor, build_year, num_room,
                  kitch_sq, state, material,
                  product_type, full_all, price_doc)

macro_s <- macro %>% select(timestamp, usdrub, unemployment)

dim(df)
```

```
## [1] 30471 13
```

```
dim(macro_s)
```

```
## [1] 2484 3
```

Converting data columns to appropriate format.

```
df$timestamp <- as.Date(df$timestamp)
macro$timestamp <- as.Date(macro$timestamp)
```

We join the data sets.

```
df <- df %>% left_join(macro_s)
dim(df)
```

```
## [1] 30471 15
```

The dataset includes 30471 observations and 292 columns.

```
split <- sample(c(rep(0, 0.75 * nrow(df)), rep(1, 0.25 * nrow(df))))
train = df[split == 0 , ]
test = df[split == 1 , ]
```

```
dim(train)
```

```
## [1] 22854 15
```

```
dim(test)
```

```
## [1] 7617 15
```

## 4 Explanatory Data Analysis

For aesthetic reasons, some outliers might have been removed from the graphs and they are not demonstrated separately. As we move forward through data, cleaning might take place as needed.

### 4.1 internal house characteristics

Here we list the house internal characteristics and analyse them

#### 4.1.1 full\_sq

Definition: total area in square meters, including loggias, balconies and other non-residential areas

Here we table the data and inspect full\_sq values. There are observations with value below 10 square meter and as they are suspicious, so we further investigate them.

```
table(train$full_sq)
```


```
##
##  0   1   5   6  12  13  14  15  16  17  18  19  20  21  22  23
##  2  19   1   1   2   7   7  11   7  13  15  19  19   8   6   6
## 24 25 26 27 28 29 30 31 32 33 34 35 36 37 38 39
## 13 19 44 75 64 198 299 404 620 234 475 441 399 882 1404 750
## 40 41 42 43 44 45 46 47 48 49 50 51 52 53 54 55
## 654 581 510 628 768 688 365 265 268 181 292 582 496 493 430 314
## 56 57 58 59 60 61 62 63 64 65 66 67 68 69 70 71
## 347 347 432 427 475 404 362 477 471 228 112 135 110 115 106 118
## 72 73 74 75 76 77 78 79 80 81 82 83 84 85 86 87
## 168 186 350 241 291 333 234 344 128 102 200 92 158 99 53 40
## 88 89 90 91 92 93 94 95 96 97 98 99 100 101 102 103
## 39 47 42 28 34 24 32 35 47 28 30 29 49 36 55 25
## 104 105 106 107 108 109 110 111 112 113 114 115 116 117 118 119
## 29 16 19 20 13 14 13 7 25 12 11 17 15 17 12 12
## 120 121 122 123 124 125 126 127 129 130 131 132 133 134 135 136
## 14 16 8 14 12 14 15 10 6 2 7 2 7 21 14 8
## 137 138 139 140 141 142 143 144 146 147 148 149 150 151 153 154
## 7 12 1 2 1 6 6 3 8 7 2 2 6 3 1 2
## 155 156 157 158 159 160 161 164 165 166 167 168 169 170 172 173
## 9 8 3 1 4 6 1 1 8 7 12 1 6 4 4 1
## 174 177 178 179 181 182 183 184 185 186 187 195 197 199 202 204
## 1 1 2 2 2 1 1 1 1 1 2 1 2 1 2 1
## 206 207 209 210 211 216 219 220 226 275 291 303 353 412 461 603
## 2 1 1 1 1 1 2 1 1 1 1 1 1 1 1 1
## 634 637 5326
## 1 1 1
```

If the area of a house is zero, we convert it to NA.

```
train[,"full_sq"][train[,"full_sq"] == 0] <- NA
```

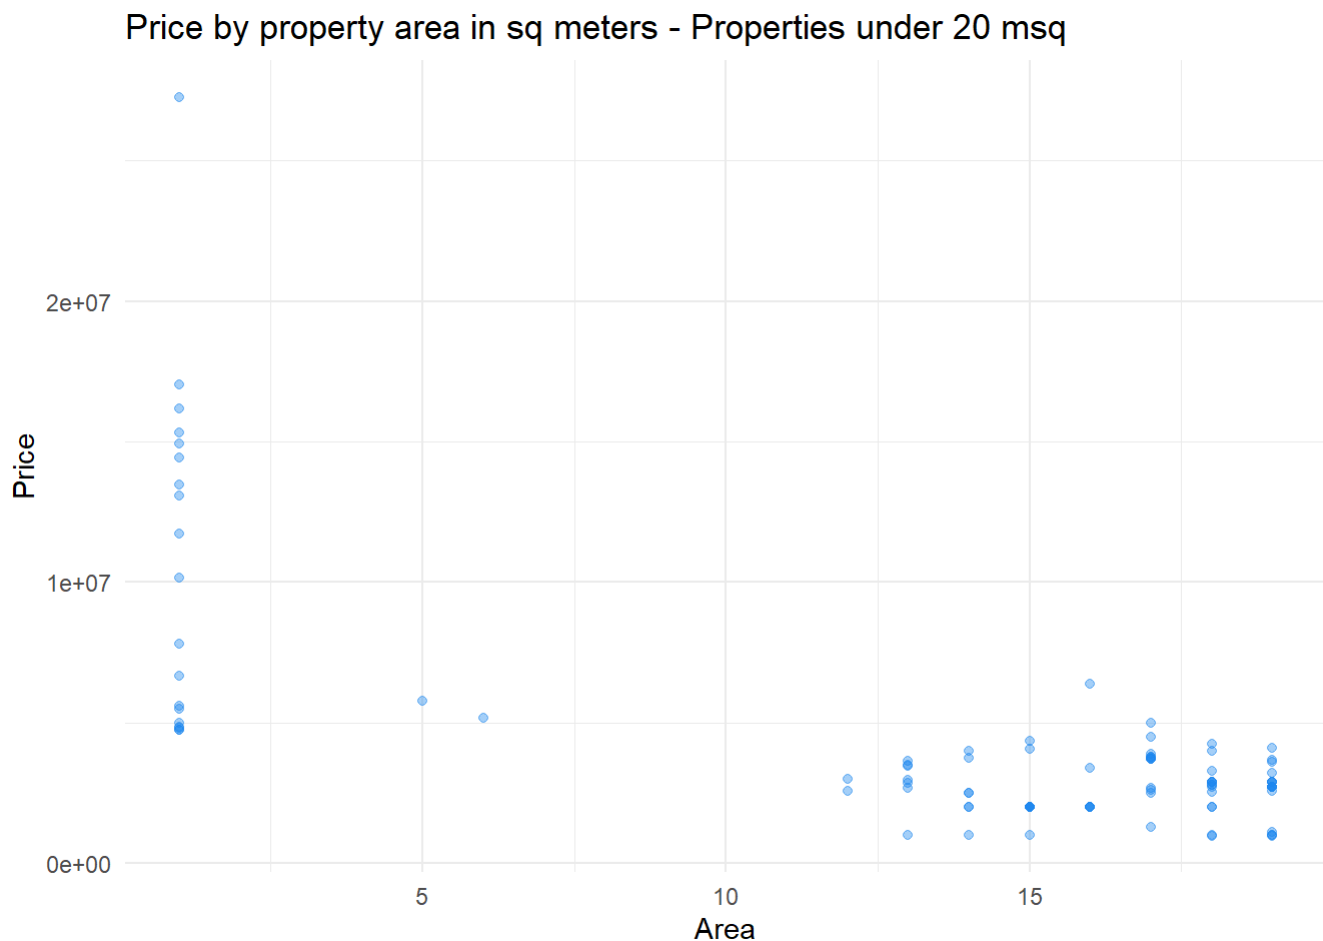
The following is a scatter plot of the price by property area.

```
train %>%
  filter(full_sq < 1000) %>%
  ggplot(aes(x=full_sq, y=price_doc)) +
  geom_point(color='dodgerblue2', alpha=0.3) +
  geom_smooth(color='deeppink4') +
  scale_y_log10() +
  labs(x='Area', y='Price', title='Price by property area in sq meters') +
  theme_minimal()
```



we graph the suspicious properties, those with an area below 20 square meter. As we are not able to further investigate the matter, we let them to stay as they are.

```
train %>%
  filter(full_sq < 20) %>%
  ggplot(aes(x=full_sq, y=price_doc)) +
  geom_point(color='dodgerblue2', alpha=0.4) +
  theme_minimal() +
  labs(x='Area', y='Price', title='Price by property area in sq meters - Properties under 20 m
sq')
```

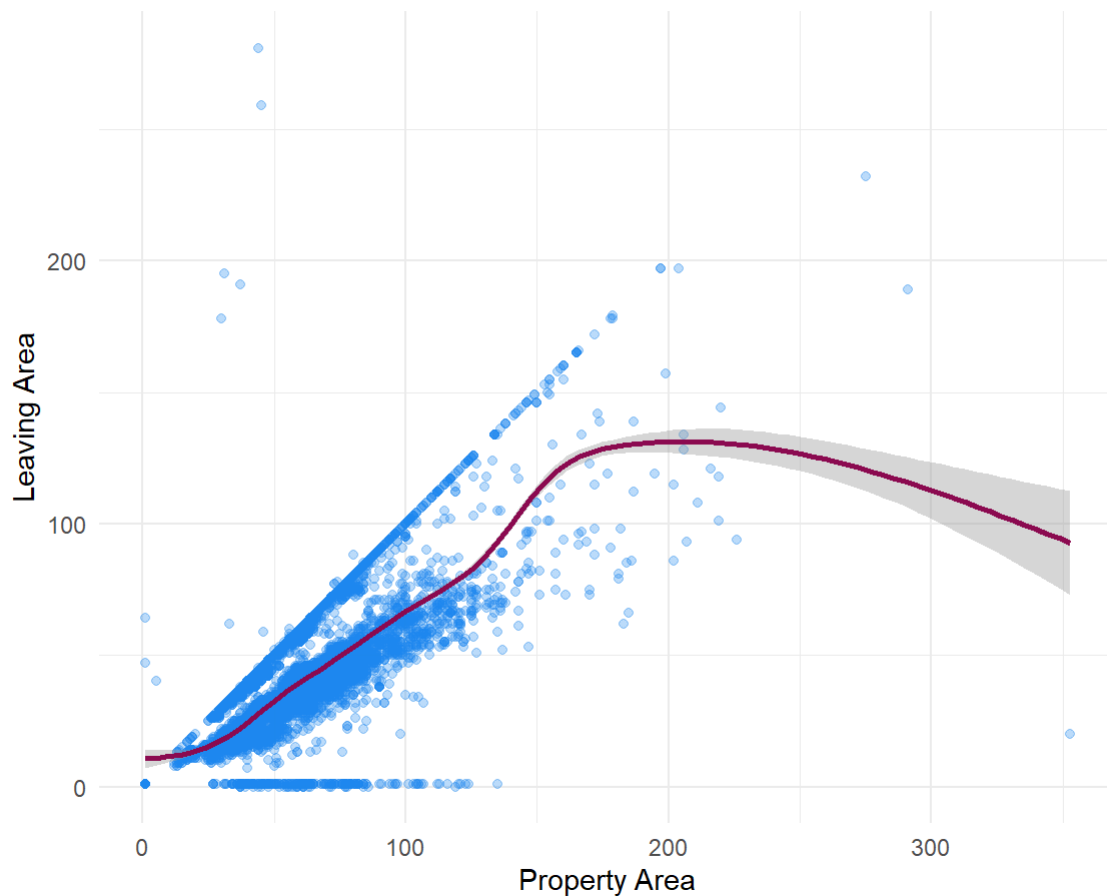


### 4.1.2 life\_sq

Next we graph leaving area against the full property area, we expect to see all values of living are below that of property area. We remove outliers from the graph to have a better view of the relation.

```
train %>%
  filter(full_sq < 400 & life_sq < 300) %>%
  ggplot(aes(y=life_sq, x=full_sq)) +
  geom_point(color='dodgerblue2', alpha=0.3) +
  geom_smooth(color = 'deeppink4') +
  coord_fixed(ratio = 1)+
  labs(y='Leaving Area' , x='Property Area',
       title='Leaving Area by Property area in sq meters') +
  theme_minimal()
```

## Leaving Area by Property area in sq meters

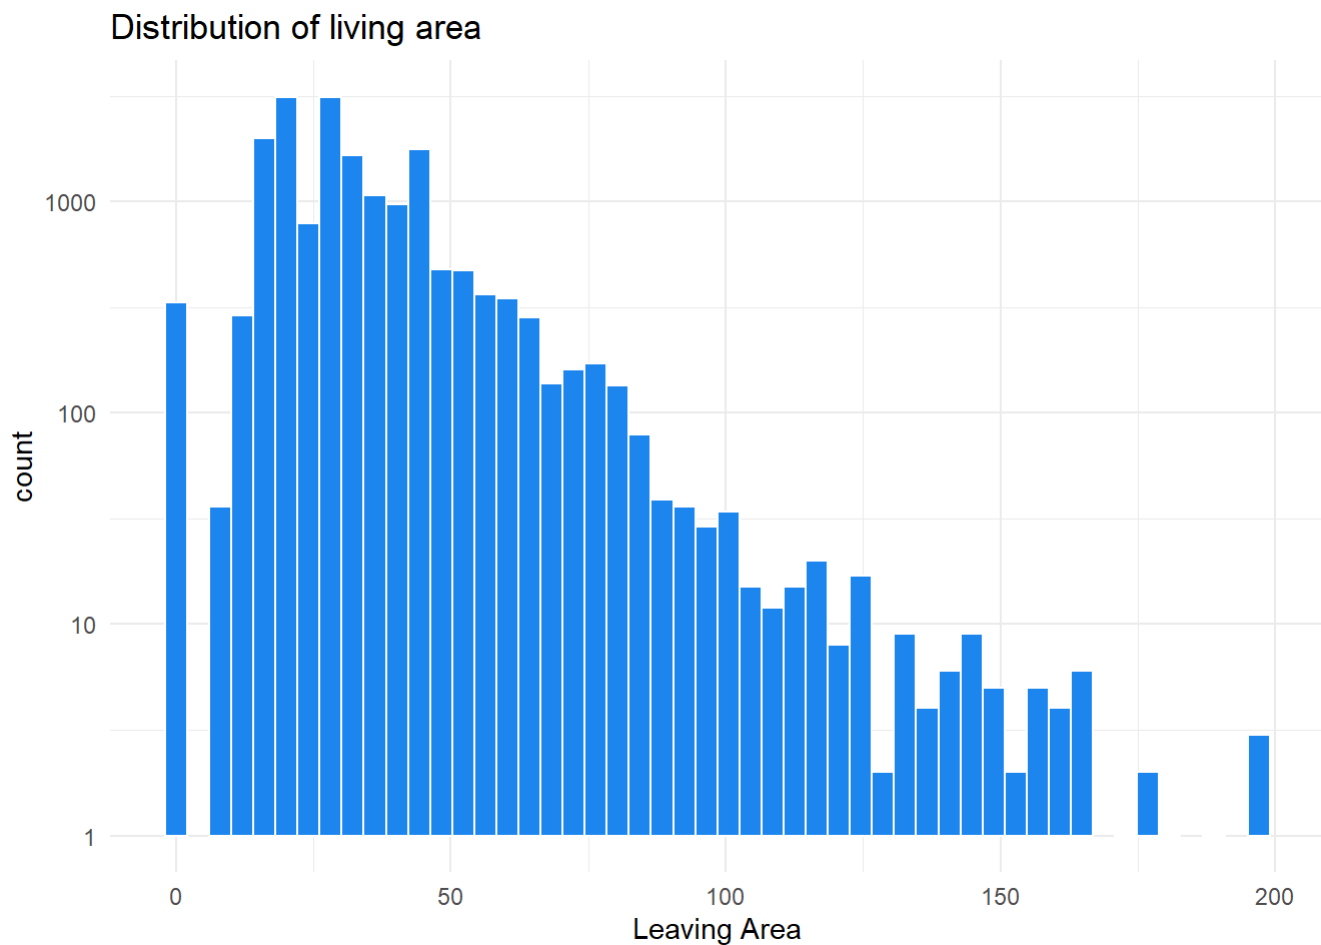


The following line of code removes the living area value of observations in which the property area is smaller than living area, as we are assuming the property value is probably more reliable.

```
train[,"life_sq"][train[,"life_sq"]>train[,"full_sq"]] <- NA
```

Now we take a look at the distribution of the leaving area.

```
train %>%
  filter(full_sq < 1000 & life_sq < 200) %>%
  ggplot(aes(x=life_sq)) +
  geom_histogram(color= "white" ,fill='dodgerblue2', bins=50) +
  scale_y_log10()+
  labs(x='Leaving Area',
       title='Distribution of living area') +
  theme_minimal()
```

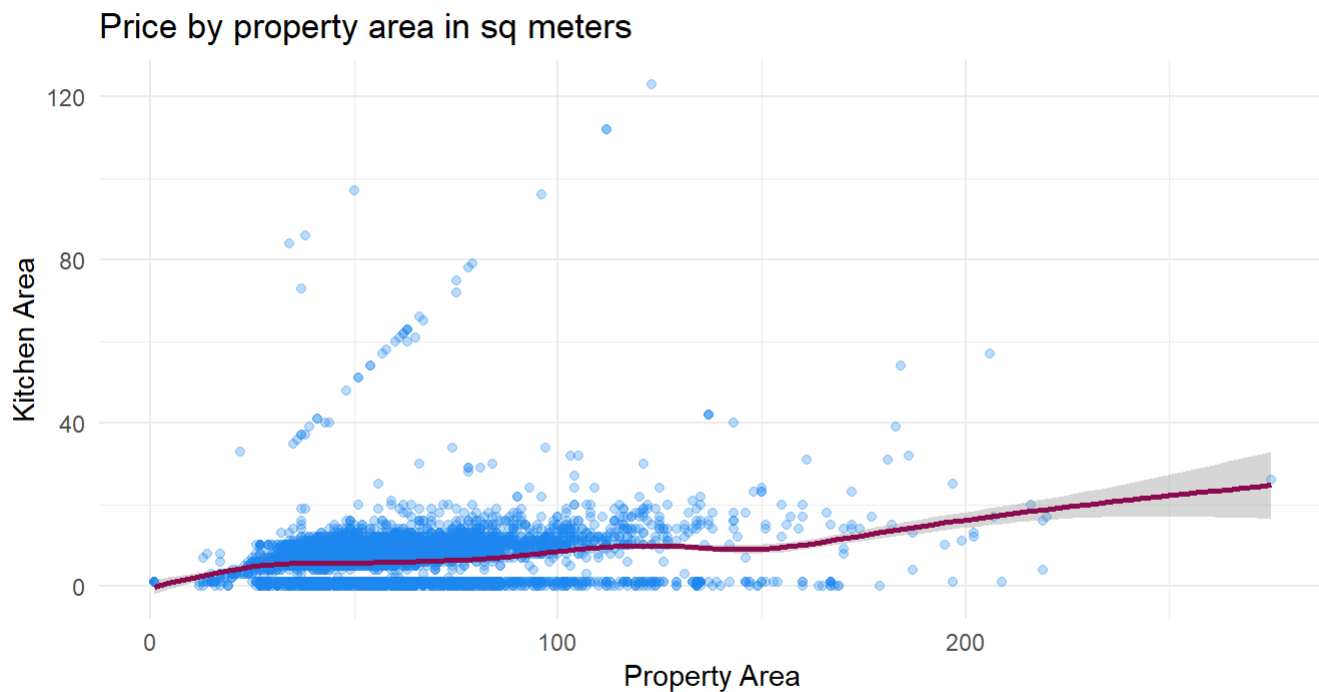


### 4.1.3 kitch\_sq

we graph the area of kitchen against the property area. As one could easily justify it, the kitchen area, increases with a small slope.

```
train %>%
  filter(full_sq < 300 & kitch_sq < 500) %>%
  ggplot(aes(y=kitch_sq, x=full_sq)) +
  geom_point(color='dodgerblue2', alpha=0.3) +
  geom_smooth(color = 'deeppink4') +
  coord_fixed(ratio = 1) +
  labs(y='Kitchen Area', x='Property Area',
       title='Price by property area in sq meters')+
  theme_minimal()
```



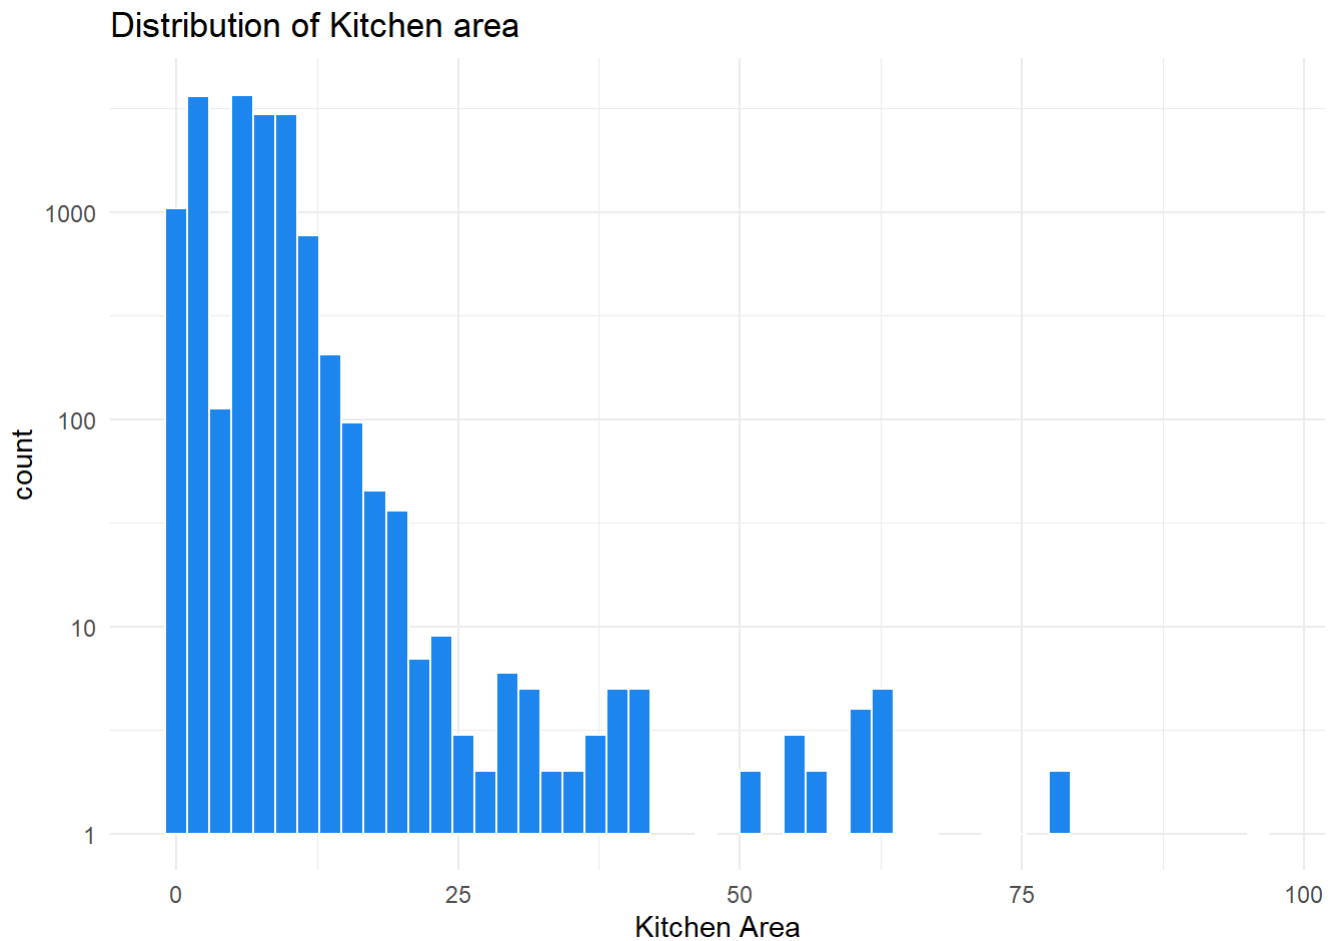


We remove kitchen values bigger than the property area.

```
train[, "kitch_sq"][train$kitch_sq > train$full_sq] <- NA
```

Here we have the histogram of kitchen area.

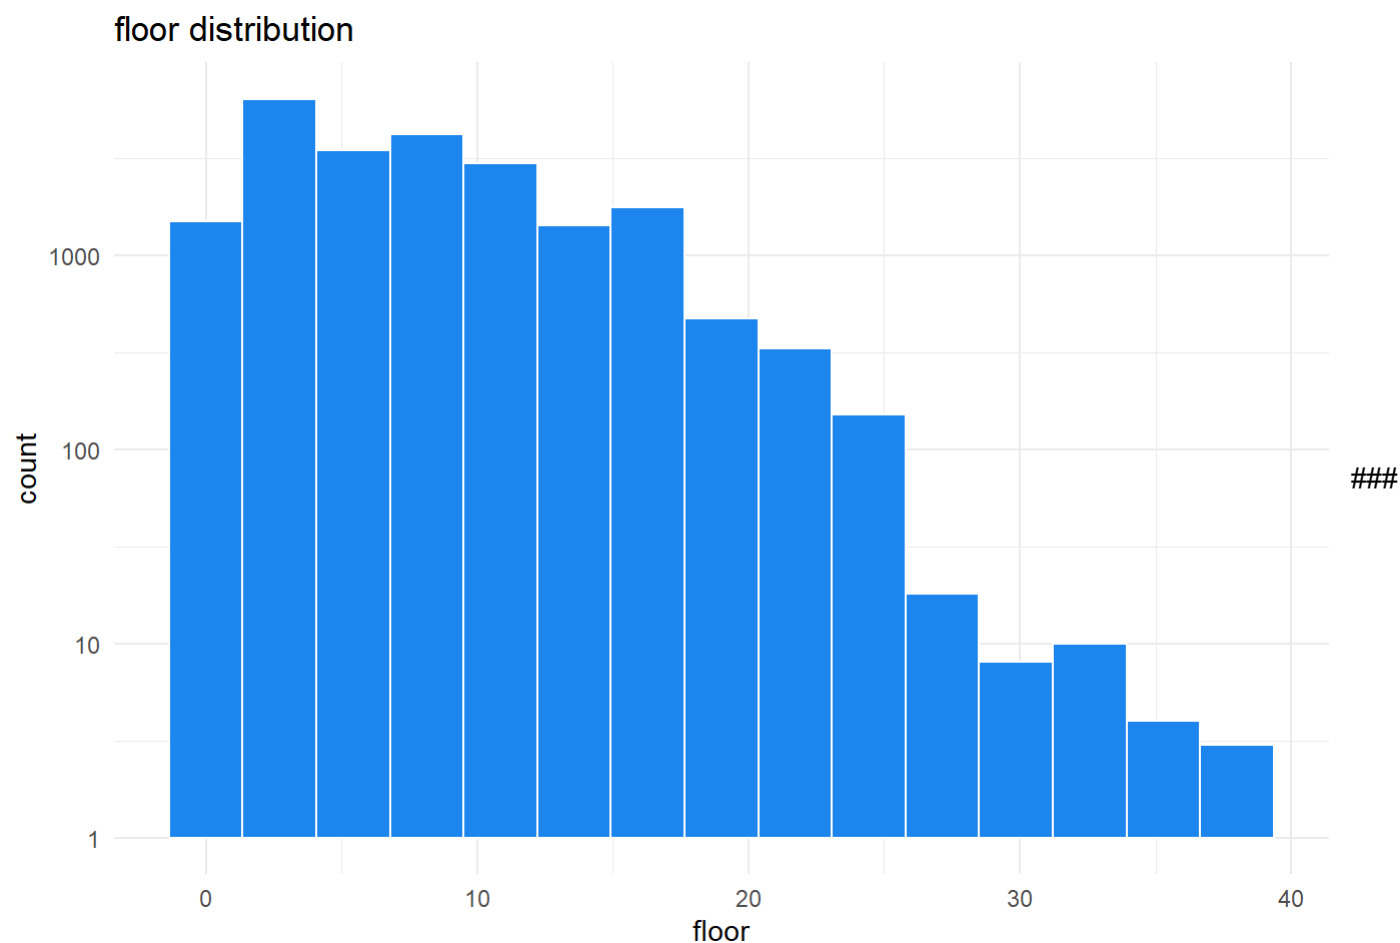
```
train %>%
  filter(kitch_sq < 100 ) %>%
  ggplot(aes(x=kitch_sq)) +
  geom_histogram(color= "white" ,fill='dodgerblue2', bins=50) +
  scale_y_log10() +
  labs(x='Kitchen Area',
       title='Distribution of Kitchen area') +
  theme_minimal()
```



## 4.1.4 floor

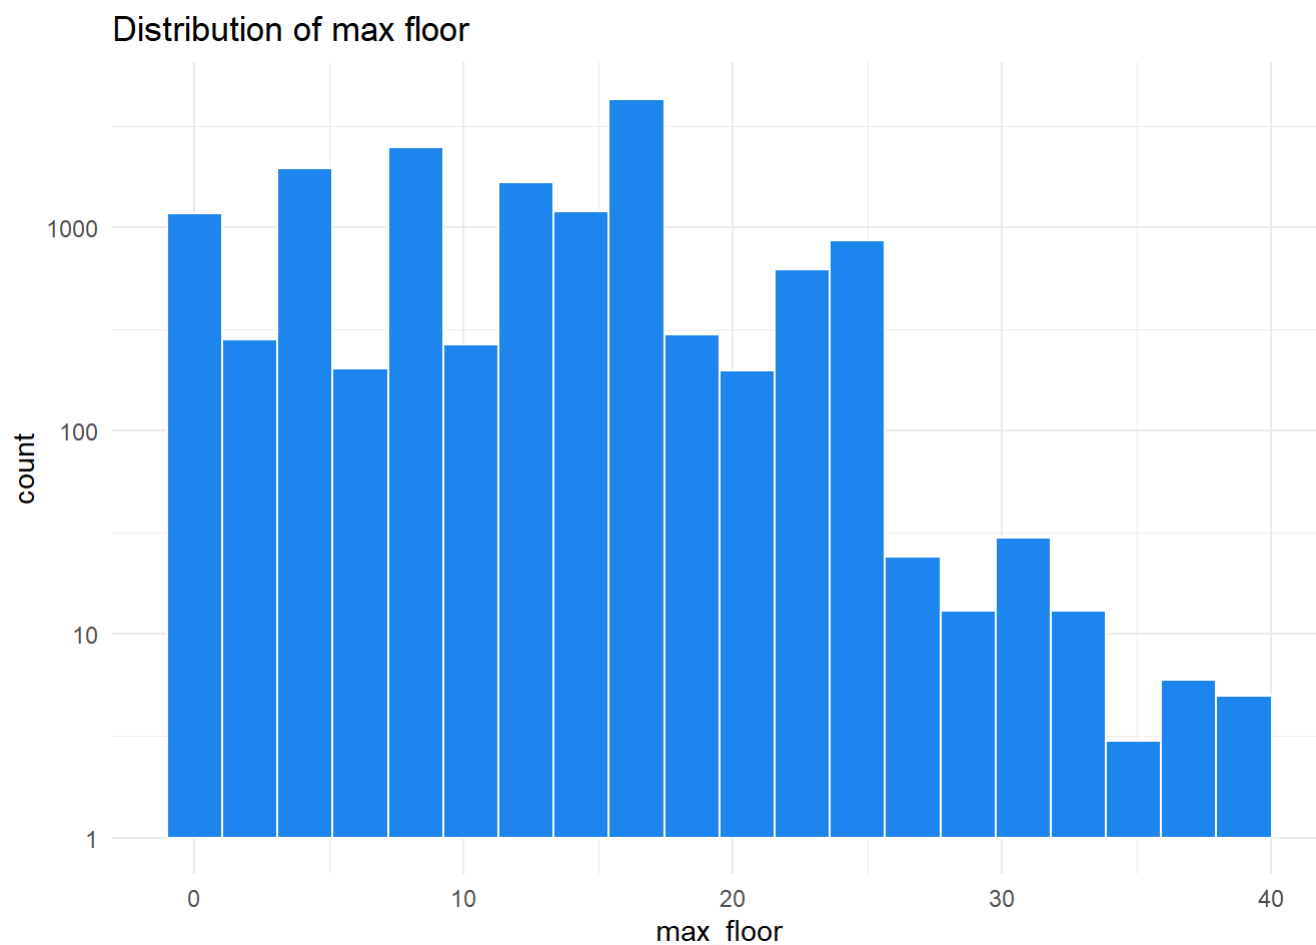
Here we have the distribution of variable floor.

```
train %>%  
  filter(floor < 40) %>%  
  ggplot(aes(x=floor)) +  
  geom_histogram(color= "white" ,fill='dodgerblue2', bins=15) +  
  scale_y_log10() +  
  labs(x='floor',  
       title='floor distribution') +  
  theme_minimal()
```




max\_floor Here the max floor

```
train %>%
  filter(max_floor < 40) %>%
  ggplot(aes(x=max_floor)) +
  geom_histogram(color= "white" ,fill='dodgerblue2', bins=20) +
  scale_y_log10() +
  ggtitle('Distribution of max floor')+
  theme_minimal()
```



We check the property floor against the maximum number of floors. we cap the graph axes on 25 floors and 25 max floors.

```
train %>%
  filter(max_floor < 25 & floor < 25) %>%
  ggplot(aes(y= floor , x= max_floor)) +
  geom_jitter(color='deeppink4', alpha=0.1) +
  coord_fixed(ratio = 1) +
  labs(x='max_floor', y='floor', title='Floor by Max Floor')+
  theme_minimal()
```




We remove max\_floors that are smaller than floors.

```
train$max_floor[train$max_floor<train$floor] <- NA
```

## 4.1.5 material

Here we table the material of the each house. We don't have list to know what the materials actually are./ There is only one observation with material 1.

```
train %>%
  ggplot( aes(x=material)) +
  geom_bar(fill = "dodgerblue1", color = "white") +
  scale_x_continuous(breaks = seq(1,6,1)) +
  geom_text(stat='count', aes(label=..count..), vjust=2)+
  theme_minimal()
```



## 4.1.6 build\_year

We first inspect the data using table command.

```
table(train$build_year)
```


```
##
##      0      1      3      71      215      1876      1886      1890
##    406    270      1      1      1      1      1      3
##   1895   1896   1900   1904   1906   1907   1910   1912
##      1      2      2      1      1      1      4      5
##   1914   1915   1917   1920   1924   1925   1926   1927
##      2      5     13      1      3      1      6      8
##   1928   1929   1930   1931   1932   1933   1934   1935
##      9     10      3      4      6      5      9      8
##   1936   1937   1938   1939   1940   1941   1943   1946
##      2     10      4      6     12      1      2      2
##   1947   1949   1950   1951   1952   1953   1954   1955
##      3      1     17     18     33     18     29     44
##   1956   1957   1958   1959   1960   1961   1962   1963
##     38     85    135    166    262    226    268    230
##   1964   1965   1966   1967   1968   1969   1970   1971
##    239    280    256    289    276    312    313    259
##   1972   1973   1974   1975   1976   1977   1978   1979
##    276    254    270    229    204    200    164    168
##   1980   1981   1982   1983   1984   1985   1986   1987
##    174    143    142    131    130    134     98    122
##   1988   1989   1990   1991   1992   1993   1994   1995
##    127    124     89     75    101     88    117    118
##   1996   1997   1998   1999   2000   2001   2002   2003
##    117    111    111     94     95    132    160    147
##   2004   2005   2006   2007   2008   2009   2010   2011
##    156    134    171    155    179    135     94    123
##   2012   2013   2014   2015   2016   2017   2018   4965
##    180    351    709    608    276    123      1      1
## 20052009
##      1
```

In main dataset we set the build years before 1860 and after 2018 to NA

```
train$build_year[train$build_year<1860 |train$build_year> 2018 ] <- NA
```

The plot of price against the built year is as follows. As it can be seen some properties values have been rounded (either by operator or sellers)

```
train %>%
  filter(build_year >1860) %>%
  ggplot(aes(y=price_doc, x=build_year)) +
  geom_point(color = 'dodgerblue2' ,alpha = .2)+
  geom_smooth(color = 'deeppink2') +
  scale_y_log10()+
  labs(x='build year', y='log Price', title='Price by build year')+
  theme_minimal()
```




Here we check the price trend in our dataset and as we see the transaction value is continuously increasing.

```
train %>%  
  ggplot(aes(y=price_doc , x= (timestamp) )) +  
  geom_smooth()+  
  scale_y_log10()+  
  labs(x='Time of transaction', y='log Price', title='Price by time of transaction')+  
  theme_minimal()
```




## Price by time of transaction



Now we check the scatter plot of price by month of transaction, to check seasonality. The transactions in spring are of a higher value compared to winter.


```
train %>%
  mutate(year = year(timestamp)) %>%
  ggplot(aes(y=price_doc , x= month(timestamp) , color = year)) +
  geom_smooth()+
  scale_y_log10()+
  scale_x_continuous(breaks = seq(1,12,1)) +
  labs(x='Month of year', y='log Price', title='Price by month of year of transaction')+
  theme_minimal()
```



## 4.1.7 num\_room

We use a histogram to investigate the number of rooms.

```
train %>%  
  ggplot(aes(x=num_room)) +  
  geom_histogram(fill = "dodgerblue2", color = "white" ,bins=20) +  
  scale_y_log10() +  
  scale_x_continuous(breaks = seq(0,10,1))
```

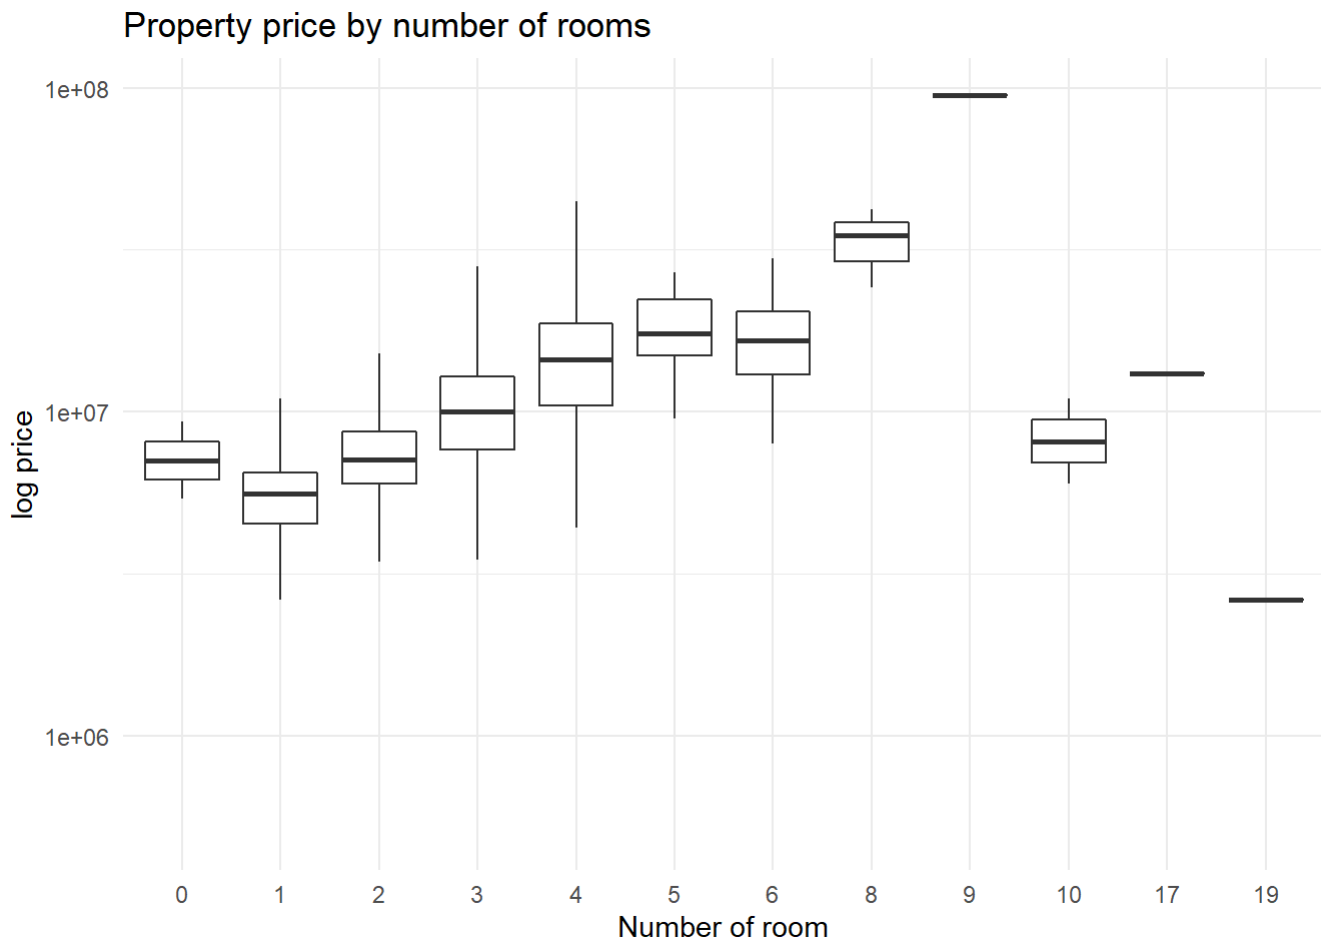


```
labs(x='Number of Rooms', y='Count', title='number of room log scaled histogram distribution')
+
  theme_minimal()
```

```
## NULL
```

We check the property price by number of rooms, as expected there is a positive correlation.


```
train %>%
  na.omit() %>%
  ggplot(aes(y=price_doc ,x=as.factor(num_room))) +
  geom_boxplot(outlier.shape = NA) +
  scale_y_log10()+
  labs(x='Number of room', y='log price', title='Property price by number of rooms')+
  theme_minimal()
```



## 4.1.8 state

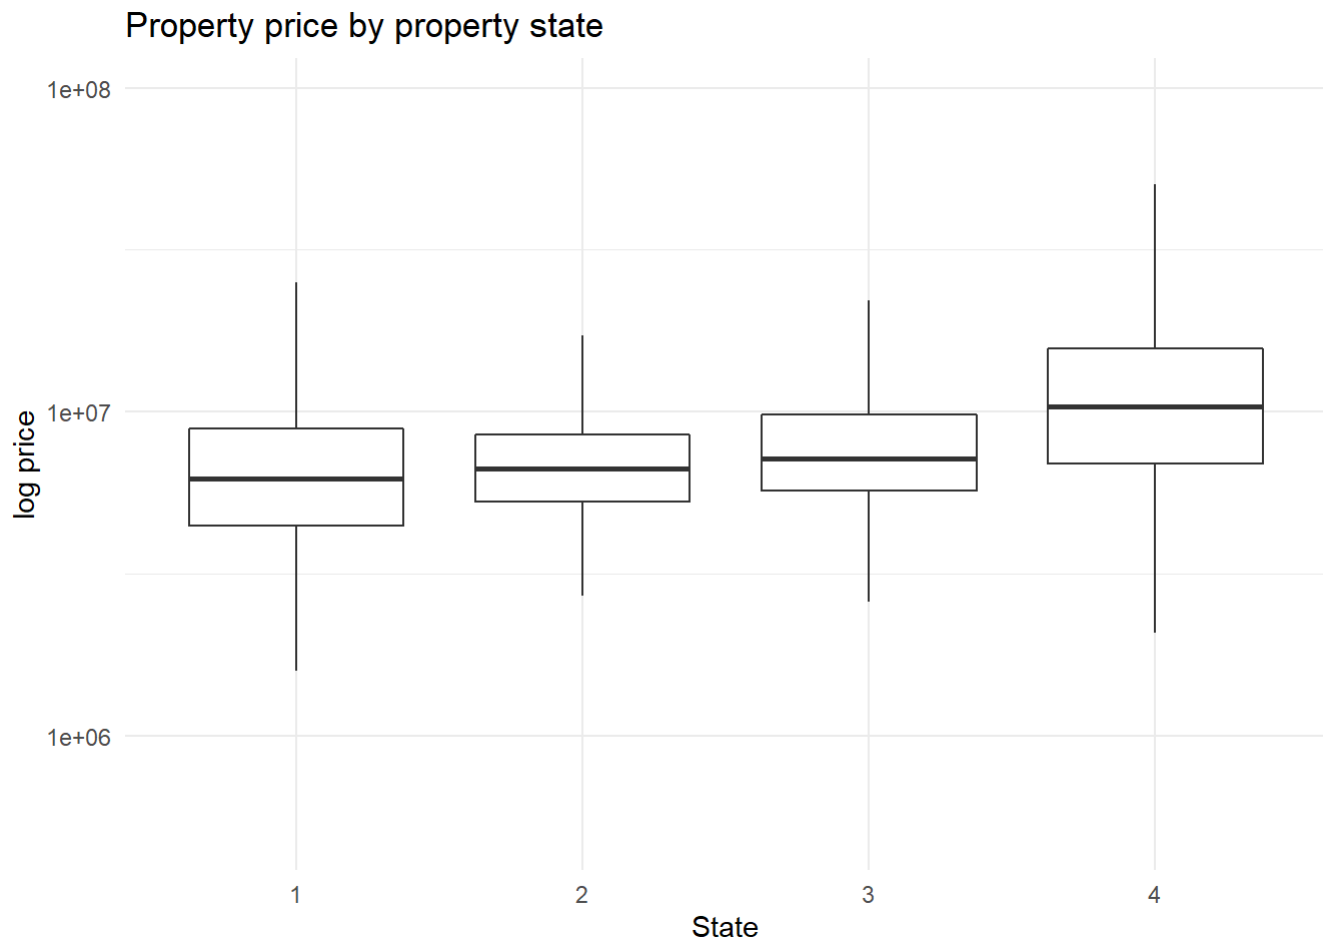
here we check the apartment condition, we also set it to factor as we don't know whether it is ordered or not. About half the data contains unknown state.

```
train$state[train$state == 33] <- 3
train %>%
  ggplot( aes(x=state)) +
  geom_bar(fill = "dodgerblue1", color = "white") +
  geom_text(stat='count', aes(label=..count..), vjust=2)+
  theme_minimal()
```



We see a slight increase in the price by state.

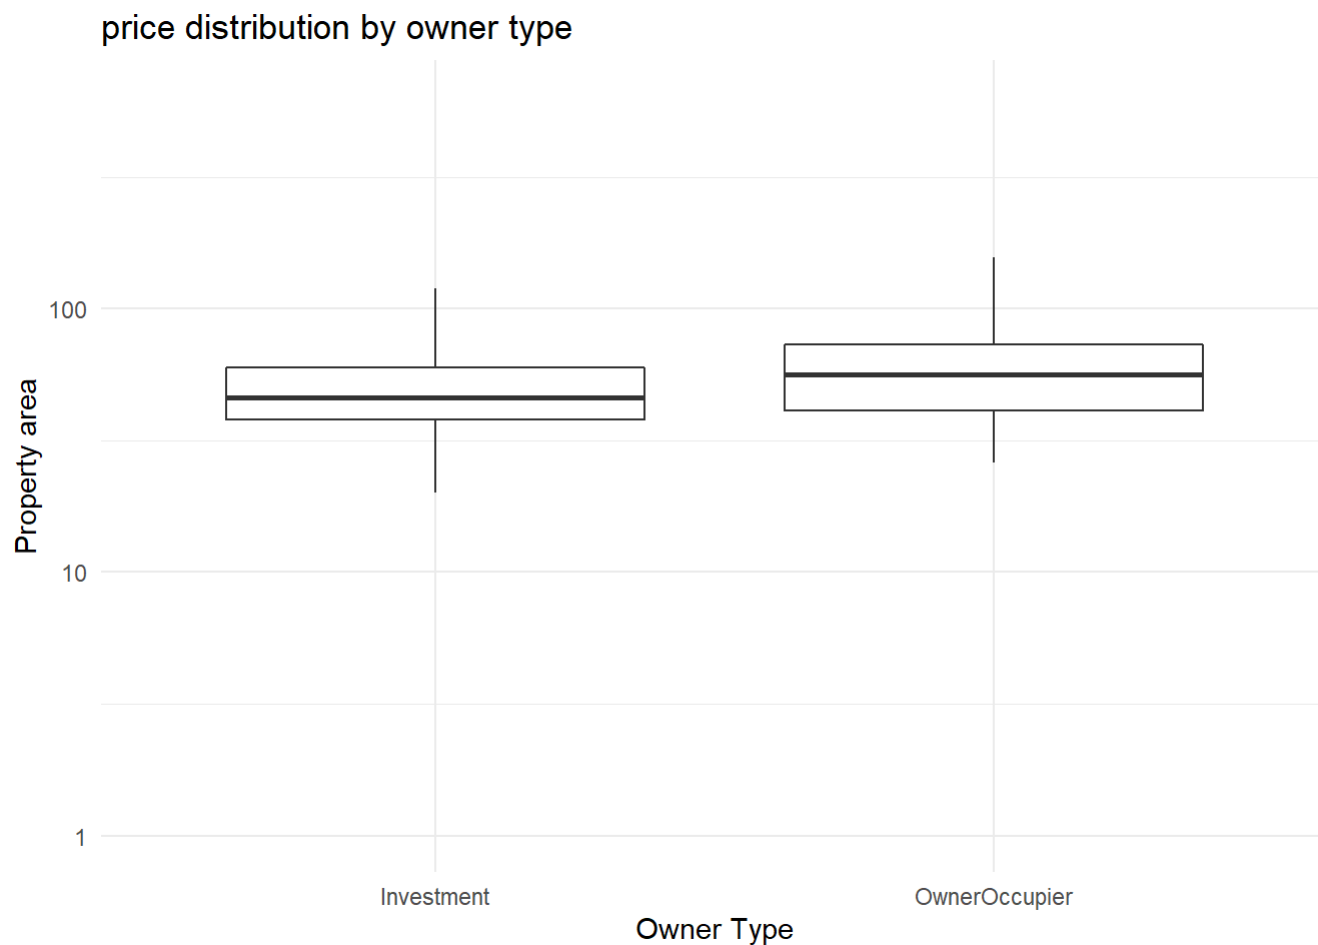
```
train %>%  
  na.omit() %>%  
  ggplot(aes(y=price_doc ,x=as.factor(state))) +  
  geom_boxplot(outlier.shape = NA) +  
  scale_y_log10()+  
  labs(x='State', y='log price', title='Property price by property state')+  
  theme_minimal()
```



### 4.1.9 product\_type

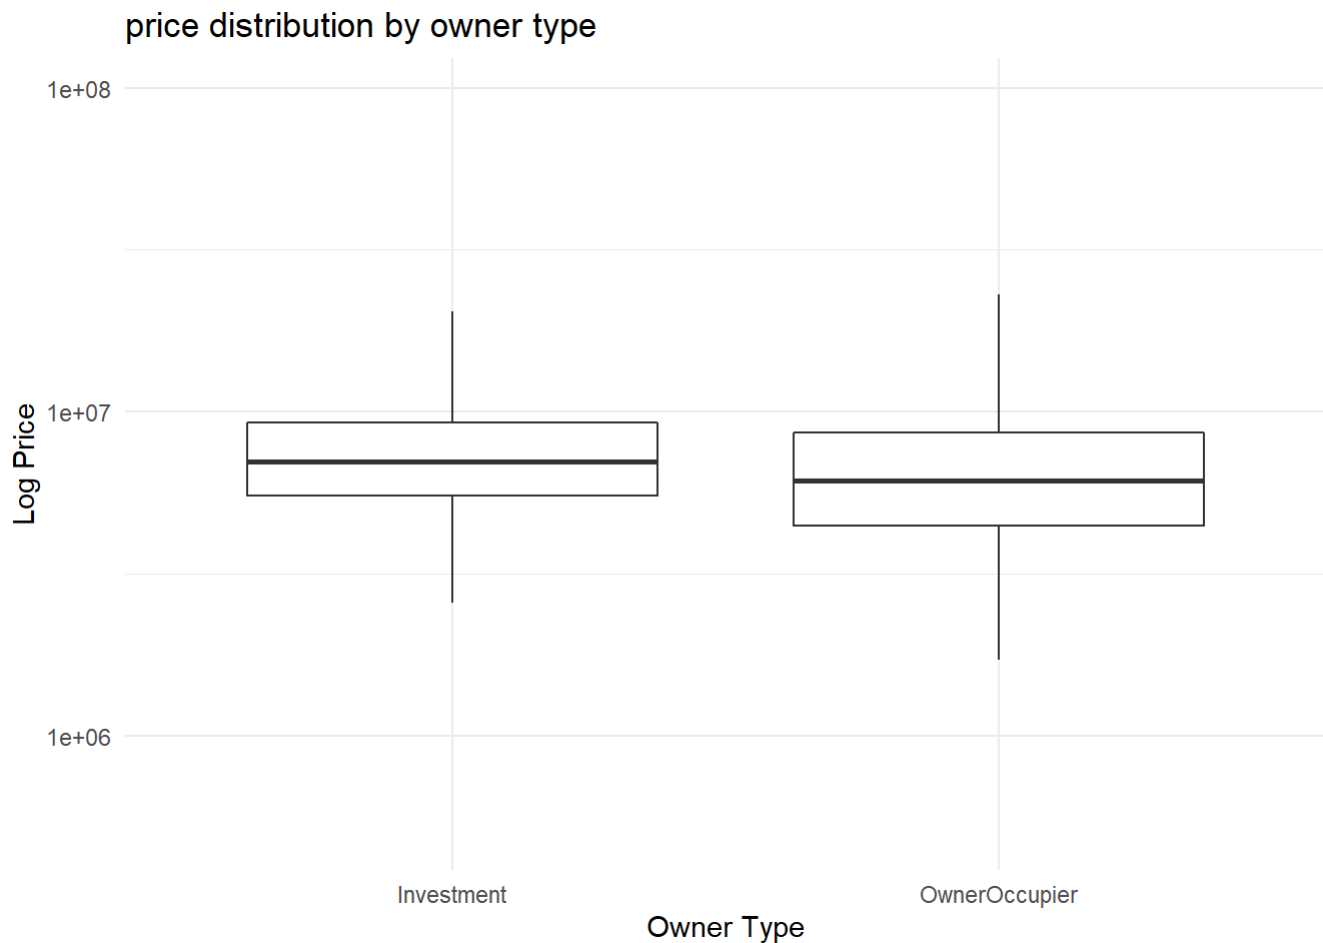
We investigate the property area against owner-occupier purchase or investment. Occupier are buying bigger houses which can be justified by the fact that they are getting both the utility of living in the property and also having it as a investment.

```
train %>%  
  na.omit() %>%  
  ggplot(aes(y=full_sq ,x=as.factor(product_type))) +  
  geom_boxplot(outlier.shape = NA) +  
  scale_y_log10()+  
  labs(x='Owner Type', y='Property area', title='price distribution by owner type')+  
  theme_minimal()
```



Here we have property value by owner against investor. Investors are buying bigger properties.

```
train %>%
  na.omit() %>%
  ggplot(aes(y=price_doc ,x=as.factor(product_type))) +
  geom_boxplot(outlier.shape = NA) +
  scale_y_log10()+
  labs(x='Owner Type', y='Log Price', title='price distribution by owner type')+
  theme_minimal()
```



## 4.2 Macro data

Among the columns of the Macro data, we have picked the most interesting ones.


### 4.2.1 usdrub

The graph is a proxy measurement of the Russia's economy. Inverting the Rubl to dollar conversion rate will give a better result, as we want to see how the value of Rubl is changing by time.

```
don <- xts(x = (1/ macro$usdrub), order.by = macro$timestamp)

dygraph(don) %>%
  dyOptions(labelsUTC = TRUE, fillGraph=TRUE, fillAlpha=0.1, drawGrid = TRUE, colors="dodgerblue
2") %>%
  dyRangeSelector() %>%
  dyCrosshair(direction = "vertical") %>%
  dyHighlight(highlightCircleSize = 5, highlightSeriesBackgroundAlpha = 0.2, hideOnMouseOut = FA
LSE) %>%
  dyRoller(rollPeriod = 1)
```






## 4.2.2 unemployment

Unemployment is another important factor

```
macro %>%
  ggplot(aes(y=unemployment , x= (timestamp) )) +
  geom_line()+
  scale_x_date(date_breaks = "years" , date_labels = "%Y") +
  labs(x='year', y='unemployment', title='Price by time of transaction')+
  theme_minimal()
```



#### 4.2.2.1 unbalanced data and sample selection

Heckman sample selection bias and unbalanced pannel data Now we left-merge the main dataset with the macro data.

now we have to clean the Test data, with the rules used on the train datasets.

```
test[,"full_sq"][test[,"full_sq"] == 0] <- NA
test[,"life_sq"][test[,"life_sq"]>test[,"full_sq"]] <- NA
test[,"kitch_sq"][test$kitch_sq>test$full_sq] <- NA
test$max_floor[test$max_floor<test$floor] <- NA
test$build_year[test$build_year<1860 |test$build_year> 2018 ] <- NA
test$state[test$state == 33] <- 3
```

## 5 Data type

We transform character vectors to factor.

```
# First we convert the train dataset characters to factor
train[sapply(train, is.character)] <- lapply(train[sapply(train, is.character)], as.factor)

# now we have to do the same for the Test, however using the factors that has been used in train
only
test$product_type <- factor(test$product_type, levels = levels(train$product_type))
sapply(train, class)
```

```
##      timestamp      full_sq      life_sq      floor      max_floor      build_year
##      "Date"        "integer"    "integer"    "integer"    "integer"    "integer"
##      num_room      kitch_sq      state      material product_type      full_all
##      "integer"    "integer"    "numeric"    "integer"    "factor"    "integer"
##      price_doc      usdrub      unemployment
##      "integer"    "numeric"    "numeric"
```

```
sapply(test, class)
```


```
##      timestamp      full_sq      life_sq      floor      max_floor      build_year
##      "Date"        "integer"    "integer"    "integer"    "integer"    "integer"
##      num_room      kitch_sq      state      material product_type      full_all
##      "integer"    "integer"    "numeric"    "integer"    "factor"    "integer"
##      price_doc      usdrub      unemployment
##      "integer"    "numeric"    "numeric"
```

## 6 Imputing the missing data

The followings are several useful links that have been used for this project. This is a book on imputation by the developer of the package mice <https://stefvanbuuren.name/fimd/ch-introduction.html> (<https://stefvanbuuren.name/fimd/ch-introduction.html>) The following is a tutorial which explains how to implement the discussed ideas. <https://amices.org/Winnipeg/> (<https://amices.org/Winnipeg/>) The following is a series of vignettes that covers the mice packages implementation. <https://www.gerkovink.com/miceVignettes/> (<https://www.gerkovink.com/miceVignettes/>)

Here we check the pattern of missing data, as we can see we have a case of multivariate missing values. In the graph, on the left we have the frequency of each pattern and on the right side the number of missing values.

```
md.pattern(train, rotate.names = TRUE)
```



##	timestamp	product_type	full_all	price_doc	usdrub	unemployment	full_sq	
## 10460	1	1	1	1	1	1	1	
## 617	1	1	1	1	1	1	1	
## 877	1	1	1	1	1	1	1	
## 199	1	1	1	1	1	1	1	
## 222	1	1	1	1	1	1	1	
## 127	1	1	1	1	1	1	1	
## 6	1	1	1	1	1	1	1	
## 48	1	1	1	1	1	1	1	
## 6	1	1	1	1	1	1	1	
## 1	1	1	1	1	1	1	1	
## 1	1	1	1	1	1	1	1	
## 5394	1	1	1	1	1	1	1	
## 227	1	1	1	1	1	1	1	
## 846	1	1	1	1	1	1	1	
## 130	1	1	1	1	1	1	1	
## 1158	1	1	1	1	1	1	1	
## 39	1	1	1	1	1	1	1	
## 159	1	1	1	1	1	1	1	
## 23	1	1	1	1	1	1	1	
## 520	1	1	1	1	1	1	1	
## 1670	1	1	1	1	1	1	1	
## 104	1	1	1	1	1	1	1	
## 18	1	1	1	1	1	1	1	
## 2	1	1	1	1	1	1	0	
##	0	0	0	0	0	0	2	
##	floor	life_sq	num_room	material	kitch_sq	max_floor	state	build_year
## 10460	1	1	1	1	1	1	1	0
## 617	1	1	1	1	1	1	1	0
## 877	1	1	1	1	1	1	0	1
## 199	1	1	1	1	1	1	0	0
## 222	1	1	1	1	1	0	1	1
## 127	1	1	1	1	1	0	1	0
## 6	1	1	1	1	1	0	0	1
## 48	1	1	1	1	1	0	0	0
## 6	1	1	1	1	0	1	1	1
## 1	1	1	1	1	0	0	1	1
## 1	1	1	1	1	0	0	1	0
## 5394	1	1	0	0	0	0	0	0
## 227	1	0	1	1	1	1	1	1
## 846	1	0	1	1	1	1	1	0
## 130	1	0	1	1	1	1	0	1
## 1158	1	0	1	1	1	1	0	0
## 39	1	0	1	1	1	0	1	1
## 159	1	0	1	1	1	0	1	0
## 23	1	0	1	1	1	0	0	1
## 520	1	0	1	1	1	0	0	0
## 1670	1	0	0	0	0	0	0	0
## 104	0	1	0	0	0	0	0	0
## 18	0	0	0	0	0	0	0	0
## 2	1	1	1	1	1	1	1	0
##	122	4790	7186	7186	7194	8332	10147	10863
								55822

Now we start the imputing the missing variables using “Multivariate Imputation by Chained Equations”.

First we set the prediction matrix.

```
pred <- imp$predictorMatrix
```

We also have to consider that the column subarea and area population have perfect correlation and we should use only one of them in our analysis. We also skip the column timestamp as it is not a numerical variable. We also won't use the column price\_doc as it is our target variable and we should not leak information.

```
pred[ , "timestamp"] <- 0  
pred[ , "full_all"] <- 0  
pred[ , "price_doc"] <- 0  
pred
```

```

##          timestamp full_sq life_sq floor max_floor build_year num_room
## timestamp          0      1      1      1          1          1      1
## full_sq            0      0      1      1          1          1      1
## life_sq            0      1      0      1          1          1      1
## floor              0      1      1      0          1          1      1
## max_floor          0      1      1      1          0          1      1
## build_year         0      1      1      1          1          0      1
## num_room           0      1      1      1          1          1      0
## kitch_sq           0      1      1      1          1          1      1
## state              0      1      1      1          1          1      1
## material           0      1      1      1          1          1      1
## product_type       0      1      1      1          1          1      1
## full_all           0      1      1      1          1          1      1
## price_doc          0      1      1      1          1          1      1
## usdrub             0      1      1      1          1          1      1
## unemployment       0      1      1      1          1          1      1
##          kitch_sq state material product_type full_all price_doc usdrub
## timestamp          1      1          1          1          0          0      1
## full_sq            1      1          1          1          0          0      1
## life_sq            1      1          1          1          0          0      1
## floor              1      1          1          1          0          0      1
## max_floor          1      1          1          1          0          0      1
## build_year         1      1          1          1          0          0      1
## num_room           1      1          1          1          0          0      1
## kitch_sq           0      1          1          1          0          0      1
## state              1      0          1          1          0          0      1
## material           1      1          0          1          0          0      1
## product_type       1      1          1          0          0          0      1
## full_all           1      1          1          1          0          0      1
## price_doc          1      1          1          1          0          0      1
## usdrub             1      1          1          1          0          0      0
## unemployment       1      1          1          1          0          0      1
##          unemployment
## timestamp          1
## full_sq            1
## life_sq            1
## floor              1
## max_floor          1
## build_year         1
## num_room           1
## kitch_sq           1
## state              1
## material           1
## product_type       1
## full_all           1
## price_doc          1
## usdrub             1
## unemployment       0

```

Now we have to set the statistical method that we want to be used for prediction of each column. The mice package makes the best choices as predictive mean matching, logistic and polynomial based on data and we have to change that for variables that we think it is necessary. The columns that do not have a missing variable do not have a method.

```
meth <- imp$meth
meth
```

```
##      timestamp      full_sq      life_sq      floor      max_floor      build_year
##           ""         "pmm"       "pmm"       "pmm"         "pmm"         "pmm"
##      num_room      kitch_sq      state      material product_type      full_all
##           "pmm"         "pmm"       "pmm"       "pmm"         ""         ""
##      price_doc      usdrub      unemployment
##           ""         ""         ""
```



Now we can run the algorithm


```
imp <- mice(train, meth = meth, pred = pred, maxit = 5 , seed = 1234 , print = FALSE)
```

We check whether there is a trend in imputation, and the data seems fine.

```
plot(imp)
```







We make a long dataframe, stacking iterations of imputation over each other, since we are using the data for prediction, it is fine to do so.

```
train_stack <- complete(imp, "long")
dim(train_stack)
```

```
## [1] 114270    17
```

Now we need to impute the test data.

```
imp1 <- mice(test, maxit=0)
```



```
pred1 <- imp1$predictorMatrix
```


```
pred1[ ,"timestamp"] <- 0
pred1[ ,"full_all"] <- 0
pred1[ ,"price_doc"] <- 0
```

```
meth1 <- imp1$meth
```

```
imp1 <- mice(test, meth = meth1, pred = pred1, maxit = 5 , seed = 1234 , print = FALSE)
```

```
plot(imp1)
```





```
test_stack <- complete(imp1, "long")
dim(test_stack)
```

```
## [1] 38085    17
```

## 7 Model Fit

For modeling we use XGBoost regressor. It is fast, and has been shown to outperform most competitors.

But first lets do a simple regression.

```
regression <- lm(price_doc ~ . , data = train_stack)
regression_pred <- predict(regression, newdata = test_stack)
reg_r2 <- sum((regression_pred - test_stack$price_doc)^2)/nrow(test_stack)
reg_r2
```

```
## [1] 1.418832e+13
```

```
train_df <- data.table(train_stack[,4:17])
test_df <- data.table(test_stack[,4:17])
train_df$product_type <- as.numeric(train_df$product_type)
test_df$product_type <- as.numeric(test_df$product_type)
```

Setting the validation dataset for XGBoost.

```
train_id <- sample(1:nrow(train_df), size = floor(0.8 * nrow(train)), replace=FALSE)
# Split in training and validation (80/20)
training <- train_df[train_id,]
validation <- train_df[-train_id,]
```

One hot encoding and setting the target variable

```
new_tr <- model.matrix(~.+0,data = training[, -c("price_doc"),with=F])
new_val<- model.matrix(~.+0,data = validation[, -c("price_doc"),with=F])
new_ts <- model.matrix(~.+0,data = test_df[, -c("price_doc"),with=F])
train_traget <- training$price_doc
val_traget <- validation$price_doc
test_target <- test_df$price_doc
```

preparing XGBoost matrix.

```
dtrain <- xgb.DMatrix(data = new_tr,label = train_traget)
dval    <- xgb.DMatrix(data = new_val,label = val_traget)
dtest   <- xgb.DMatrix(data = new_ts,label = test_target)
```

Setting default default parameters for the first run.

```
params <- list(booster = "gbtree", objective = "reg:squarederror",
               eta=0.3, gamma=0, max_depth=6, min_child_weight=1,
               subsample=1, colsample_bytree=1)
```

Running the first run

```
set.seed(1234)
xgb_base <- xgb.train (params = params,
                      data = dtrain,
                      nrounds =1000,
                      print_every_n = 200,
                      eval_metric = 'rmse',
                      early_stopping_rounds = 50,
                      watchlist = list(train= dtrain, val= dval))
```

Now we run a random parameter search with 1000 iteration

```
# strt time
start.time <- Sys.time()

# empty lists
lowest_error_list = list()
parameters_list = list()

# 1000 rows with random hyperparameters
set.seed(1234)
for (iter in 1:1000){
  param <- list(booster = "gbtree",
               objective = "reg:squarederror",
               max_depth = sample(3:10, 1),
               eta = runif(1, .01, .3),
               subsample = runif(1, .7, 1),
               colsample_bytree = runif(1, .6, 1),
               min_child_weight = sample(0:10, 1)
  )
  parameters <- as.data.frame(param)
  parameters_list[[iter]] <- parameters
}

# object that contains all randomly created hyperparameters
parameters_df = do.call(rbind, parameters_list)

# using randomly created parameters to create 1000 XGBoost-models
for (row in 1:nrow(parameters_df)){
  set.seed(20)
  mdcv <- xgb.train(data=dtrain,
                   booster = "gbtree",
                   objective = "reg:squarederror",
                   max_depth = parameters_df$max_depth[row],
                   eta = parameters_df$eta[row],
                   subsample = parameters_df$subsample[row],
                   colsample_bytree = parameters_df$colsample_bytree[row],
                   min_child_weight = parameters_df$min_child_weight[row],
                   nrounds= 300,
                   eval_metric = "rmse",
                   early_stopping_rounds= 30,
                   watchlist = list(train= dtrain, val= dval)
  )
  lowest_error <- as.data.frame(1 - min(mdcv$evaluation_log$val_error))
  lowest_error_list[[row]] <- lowest_error
}

# object that contains all accuracy's
lowest_error_df = do.call(rbind, lowest_error_list)

# binding columns of accuracy values and random hyperparameter values
randomsearch = cbind(lowest_error_df, parameters_df)

# end time
```

```
end.time <- Sys.time()
time.taken <- end.time - start.time
time.taken
```

```
time.taken
```

```
## Time difference of 1.790259 hours
```

Here we have a table of our random search results

```
randomsearch <- as.data.frame(randomsearch) %>%
  rename(val_acc = `1 - min(mdcv$evaluation_log$val_error)`) %>%
  arrange(-val_acc)
```

We calculate the error of the best model on the validation set.

```
# Tuned-XGBoost model
set.seed(1234)
params <- list(booster = "gbtree",
               objective = "reg:squarederror",
               max_depth = randomsearch[1,]$max_depth,
               eta = randomsearch[1,]$eta,
               subsample = randomsearch[1,]$subsample,
               colsample_bytree = randomsearch[1,]$colsample_bytree,
               min_child_weight = randomsearch[1,]$min_child_weight)
xgb_tuned <- xgb.train(params = params,
                      data = dtrain,
                      nrounds = 1000,
                      print_every_n = 100,
                      eval_metric = "rmse",
                      early_stopping_rounds = 30,
                      watchlist = list(train= dtrain, val= dval))

# Make prediction on dvalid
validation$pred_survived_tuned <- predict(xgb_tuned, dval)

val_r2 = sum((validation$price_doc - validation$pred_survived_tuned) ^ 2) / nrow(validation)
val_r2
```

```
val_r2
```

```
## [1] 5.080351e+12
```

And finally here we have error on the test set.



```
set.seed(1234)
params <- list(booster = "gbtree",
               objective = "reg:squarederror",
               max_depth = randomsearch[1,]$max_depth,
               eta = randomsearch[1,]$eta,
               subsample = randomsearch[1,]$subsample,
               colsample_bytree = randomsearch[1,]$colsample_bytree,
               min_child_weight = randomsearch[1,]$min_child_weight)
xgb_tuned <- xgb.train(params = params,
                      data = dtrain,
                      nrounds = 1000,
                      eval_metric = "rmse",
                      early_stopping_rounds = 30,
                      watchlist = list(train= dtrain, val= dtest))
# Make prediction on dvalid
test_df$pred_price_tuned <- predict(xgb_tuned, dtest)

test_r2 = sum((test_df$price_doc - test_df$pred_price_tuned) ^ 2 ) / nrow(test_df)
test_r2
```

```
test_r2
```

```
## [1] 7.966814e+12
```

As one would expect, a randomly tuned XGBoost, drastically outperforms simple regression

```
round(test_r2/reg_r2,2)
```

```
## [1] 0.56
```